

# Experimental Design & A/B Testing

MKTG 6279

Week 1

# This course

- ▶ Basics of causal inference
- ▶ Customer segmentation
- ▶ Feature selection
- ▶ Recommendation engines

## This week

- ▶ Experiments
- ▶ A/B testing
- ▶ Causal forests
- ▶ Power tests

# Experiments

## What is an experiment?

- ▶ Take sample from a population of subjects
- ▶ Randomly assign subjects to either treatment or control
- ▶ Measure average outcome for each subjects
- ▶ Compare outcomes between treatment and control groups

This procedure allows you to estimate the causal effect of the **treatment** on some **outcome**

## Why experiments?

- ▶ Causation vs. correlation: experiments help establish cause-and-effect relationships
- ▶ Reduce risk: test changes on a small scale before rolling them out to everyone
- ▶ Optimize ROI: identify what works best and allocate resources effectively

Example: testing different ad creatives to see which generates the most clicks

## Randomized Control Trials (RCTs): The Gold Standard

Randomization: participants are randomly assigned to one of the treatment groups or a control group

- ▶ Control Group: receives the existing treatment or no treatment
- ▶ Treatment Group: receives the new treatment or intervention
- ▶ Minimizes bias (incorrect estimate) and allows for *causal* inference

# Potential Outcomes in Causal Inference

Each subject has two potential outcomes in the experiment

- ▶  $Y_i(1)$ : outcome from being treated
- ▶  $Y_i(0)$ : outcome from not being treated

The individual treatment effect is the difference between these two outcomes:

$$\tau_i = Y_i(1) - Y_i(0)$$

Unfortunately, we can only observe one potential outcome for each subject, so we can never know what  $\tau_i$  is for each subject

## Average Treatment Effects

Since we can't observe these individual effects, we start with the **Average Treatment Effect (ATE)**

$$\tau = \mathbb{E}_i[Y_i(1) - Y_i(0)]$$

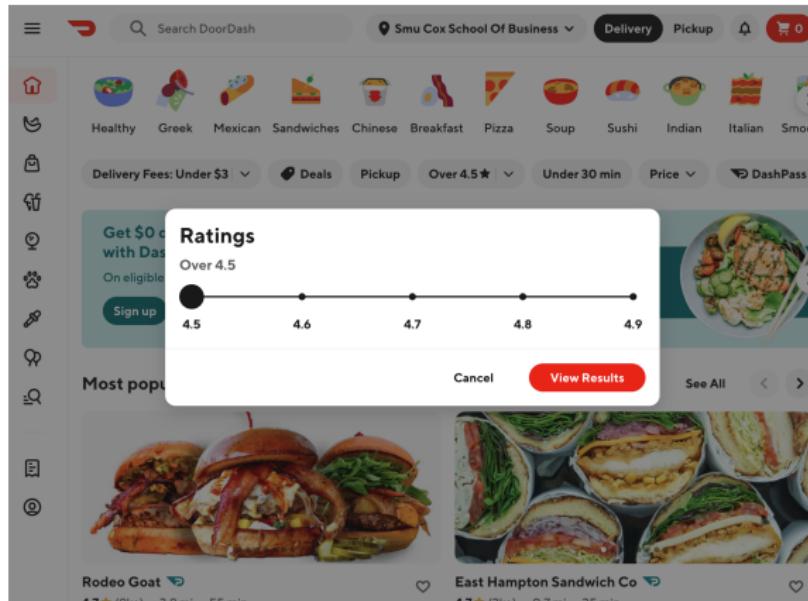
which we estimate as the difference-in-means

$$\hat{\tau} = \frac{1}{n_1} \sum Y_i(1) - \frac{1}{n_0} \sum Y_i(0)$$

where  $n_1$  and  $n_0$  are the number of treated and control subjects.

# Simple Example: DoorDash Ratings Scale

DoorDash changes the ratings minimum from 4.5 to 4



## Why might DoorDash test this?

A lower minimum range expands the range of options

But, this increases search costs and can lead to abandoned carts since “worse” restaurants are included in the results list

Use a RCT to see what happens...

## Setup

**Control:** current design (4.5+)

**Treatment:** new, lower range (4+)

**Outcome:** time on site (in seconds)

Other potential outcomes: purchase, amount spent, rate of cart abandonment, etc.

- ▶ The outcome selected depends on who cares about these results and what the goal of the analysis is

## Doordash Data

```
load('data/doordash.rdata')
str(doordash)
```

```
## 'data.frame':      1000 obs. of  7 variables:
##   $ ip        : num  10001 10002 10003 10004 10005 ...
##   $ age       : num  59 63 59 39 27 22 45 21 43 40 ...
##   $ mthsActive: num  19.26 11.51 4.26 9.53 19.54 ...
##   $ treatment : num  0 0 1 1 0 1 1 1 1 0 ...
##   $ mobile    : num  1 1 0 1 0 0 1 1 1 1 ...
##   $ seconds   : num  116.3 146.6 159.8 120.9 40.3 ...
##   $ purchase  : int  0 0 1 0 0 0 1 0 0 0 ...
```

# ATE

```
ate = doordash %>%
  group_by(treatment) %>%
  summarise(seconds_bar = mean(seconds))
ate
```

```
## # A tibble: 2 x 2
##   treatment seconds_bar
##       <dbl>      <dbl>
## 1          0      99.3
## 2          1     105.
```

So the ATE is about 6.1 seconds.

## Another way...

```
with(doordash,  
     mean(seconds[treatment == 1]) -  
     mean(seconds[treatment == 0]))  
  
## [1] 6.051932
```

Is this significantly different from zero?

## Use a t-test

Testing the difference in means (i.e., average seconds on the site) between the **treated** and **control** groups

```
t.test(seconds ~ treatment, doordash)
```

```
##  
##  Welch Two Sample t-test  
##  
## data: seconds by treatment  
## t = -2.6885, df = 982.28, p-value = 0.0073  
## alternative hypothesis: true difference in means between  
## 95 percent confidence interval:  
## -10.469417 -1.634448  
## sample estimates:  
## mean in group 0 mean in group 1  
##           99.25854          105.31048
```

Or use regression...

Slightly different 95% CI on treatment, but this is splitting hairs...

```
r1 = lm(seconds ~ treatment, doordash)
tidy(r1) %>% select(1,2,5) %>%
  mutate(across(-1, ~round(., 3)))
```

```
## # A tibble: 2 x 3
##   term      estimate p.value
##   <chr>     <dbl>    <dbl>
## 1 (Intercept) 99.3     0
## 2 treatment    6.05    0.007
```

```
confint(r1)
```

```
##                   2.5 %     97.5 %
## (Intercept) 96.203769 102.31332
## treatment    1.642745 10.46112
```

## Why use experiments?

The difference-in-means  $\hat{\tau}$  is an **unbiased** estimate of the average treatment effect (ATE)

- ▶ If we only gave mobile users the treatment and desktop users the control, our estimate may be biased because it captures both the treatment effect *and* the effect of the device type

In other words, randomization eliminates the possibility of **confounders**

## SUTVA

An assumption is that the outcomes for one subject do not affect the outcomes for other subjects

This is called the **Stable Unit Treatment Value Assumption** (SUTVA)

- ▶ Usually a problem when subjects might share their experiences with others (e.g., if other subjects become aware that a test is occurring)

## Aside: random sampling

The average treatment effect averages over the **population in the experiment**.

The ideal experiment randomly samples from the target population.

- ▶ Trivial in engineering experiments (e.g., Amazon, DoorDash, etc.)
- ▶ Do-able in marketing experiments
- ▶ Unethical in clinical trials

So, there are two types of randomization in the ideal experiment: **random sampling** to select subjects and **random assignment** to treatments

In later classes we talk about how to deal with violations of these.

## Wait, what if the treatment effect depends on...

- ▶ City vs rural individuals?
- ▶ Mobile vs desktop users?
- ▶ DashPass member status?
- ▶ Age?
- ▶ Tenure with DoorDash?
- ▶ Anything else. . .

## Enter HTE and CATE

When treatment effects vary by subgroup (even at the individual level) we have **heterogeneous treatment effects** (HTEs)

We allow for HTEs by estimating **conditional average treatment effects** (CATEs):

$$\mathbb{E}[\tau_i | X] = \mathbb{E}[Y_i(1) - Y_i(0) | X_i]$$

“The estimated treatment effect *conditional on* some value  $X_i$ ”

## DoorDash CATE

Does the treatment effect *depend on* whether this was a mobile (vs desktop) user?

```
r2=lm(seconds ~ treatment*mobile, doordash)
tidy(r2) %>% select(1,2,5) %>%
  mutate(across(-1, ~round(., 3)))
```

```
## # A tibble: 4 x 3
##   term           estimate p.value
##   <chr>         <dbl>    <dbl>
## 1 (Intercept)  94.4     0
## 2 treatment    3.73    0.226
## 3 mobile       10.1    0.001
## 4 treatment:mobile 4.72    0.287
```

## Confidence intervals

```
confint(r2)
```

```
##                                2.5 %    97.5 %
## (Intercept)      90.246986 98.583938
## treatment        -2.304148  9.756347
## mobile           4.090144 16.137987
## treatment:mobile -3.975074 13.411416
```

## DoorDash CATE

But make sure to control for potentially confounding variables for more precise estimates...

```
r3=lm(seconds ~ treatment*mobile + age + mthsActive,  
      doordash)  
coef(r3)
```

##	(Intercept)	treatment	mobile
##	45.215676	7.897196	11.860501
##	mthsActive	treatment:mobile	
##	-2.996878	6.104693	

## More precise CIs

```
confint(r3)
```

```
##                      2.5 %    97.5 %
## (Intercept)      44.961696 45.469656
## treatment        7.719270  8.075123
## mobile           11.682812 12.038191
## age              1.991776  2.001127
## mthsActive       -3.005754 -2.988002
## treatment:mobile 5.848315  6.361070
```

## Use Causal Forests for more flexible mappings...

Rather than use linear regression to define subgroups, use random forests to accommodate HTEs

Instead of predicting outcomes, predict treatment effects

- ▶ *Random* forests split data to minimize prediction errors
- ▶ *Causal* forests split data to maximize the differences in estimated treatment effects between subgroups

## “Honest” tree construction

In a causal forest, data is usually split into separate subsets:

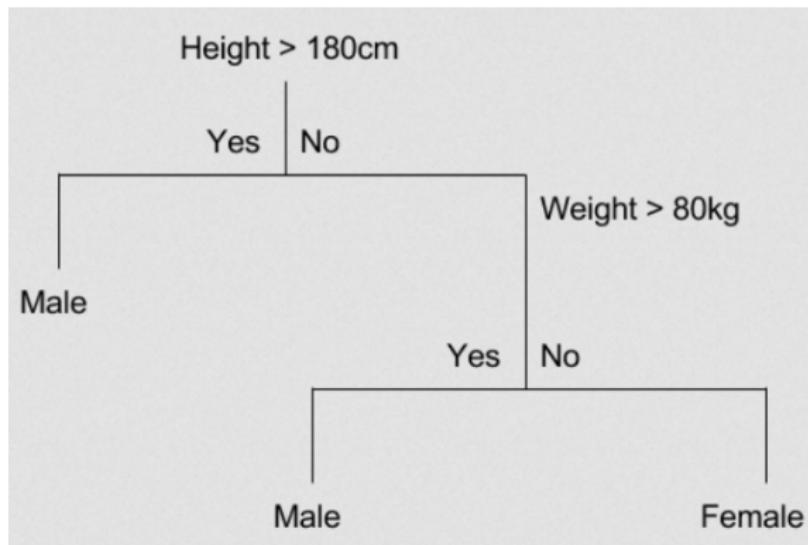
- 1) One for building the tree structure
- 2) Another for estimating treatment effects within the resulting nodes

Also known as “split-sample” or “double-sample” estimation

## Driven by CART

A **classification and regression tree (CART)** is the building block for causal forests

These partition the data into subsets based on feature values



# Classification tree for categorical outcomes

## Example

- ▶ Predicting churn (yes/no)

## Splitting criteria

- ▶ Gini Impurity: probability of incorrectly classifying a randomly chosen element
- ▶ Entropy: disorder or randomness in a set of data
- ▶ Minimize impurity in the child nodes
- ▶ Intuition: a group of {A,A,A} has less impurity in the final node versus {A,B,C}, so create branches and nodes accordingly

## Prediction

- ▶ Most common value within a node

# Regression tree for continuous outcomes

## Example

- ▶ Predicting house prices

## Splitting criteria

- ▶ Minimize the sum of squared residuals (differences between predicted and actual values) within the child nodes
- ▶ In other words, create groups of similar values

## Prediction

- ▶ Average value of outcomes within a node

## Simple regression tree in DoorDash

We can fit a random forest using the grf package

grf stands for “Generalized Random Forests”

```
library(grf)
X = doordash %>% select(mobile,age,mthsActive)
Y = doordash$seconds

rf    = regression_forest(X, Y, num.trees = 1)
tree = get_tree(rf, 1)
#plot(tree)
```

## Random forests

A **random forest** is a collection of trees estimated from different sub-samples of the data

The predictions for each unit are averaging the predictions across trees

This is an example of **bagging** which is an **ensemble technique**

Simply increase `num.trees` from the prior slide to turn the single “tree” into a “forest”

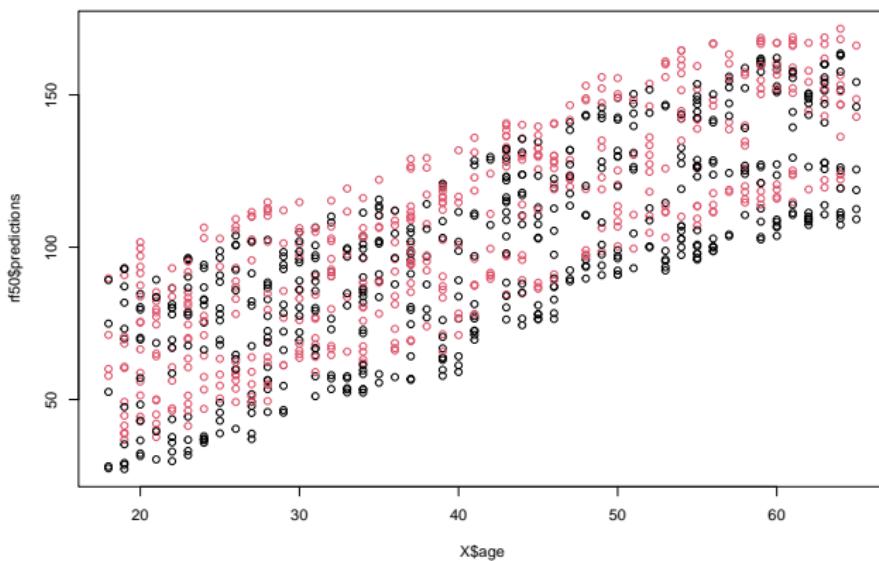
```
rf50 = regression_forest(X, Y, num.trees = 50)
#plot(get_tree(rf50, 1))
#plot(get_tree(rf50, 50))
```

## Random forest predictions

Random forests aggregate across trees to get the prediction

Unlike linear regression, predictions don't *necessarily* follow a linear pattern

```
plot(X$age, rf50$predictions, col=factor(X$mobile))
```



## Now for *causal* forests

A **causal forest** is a random forest built to predict each unit's *treatment effect*  $\tau_i = Y_i(1) - Y_i(0)$  as a function of potential variables  $X_i$

It allows us to sort through the potential variables quickly and identify non-linear relationships

## Causal forest for DoorDash

To fit a causal forest, we call `grf::causal_forest`

The inputs are the variables X, the outcomes Y, the treatment W, and the probability of treatment in the experiment W.hat

```
W      = doordash$treatment  
W.hat = .5 #randomly assigned, equally sized groups  
cf = causal_forest(X, Y, W, W.hat=0.5, seed=1)
```

## Causal forest CATEs

The goal of the causal forest is to estimate the conditional average treatment effects (CATEs):

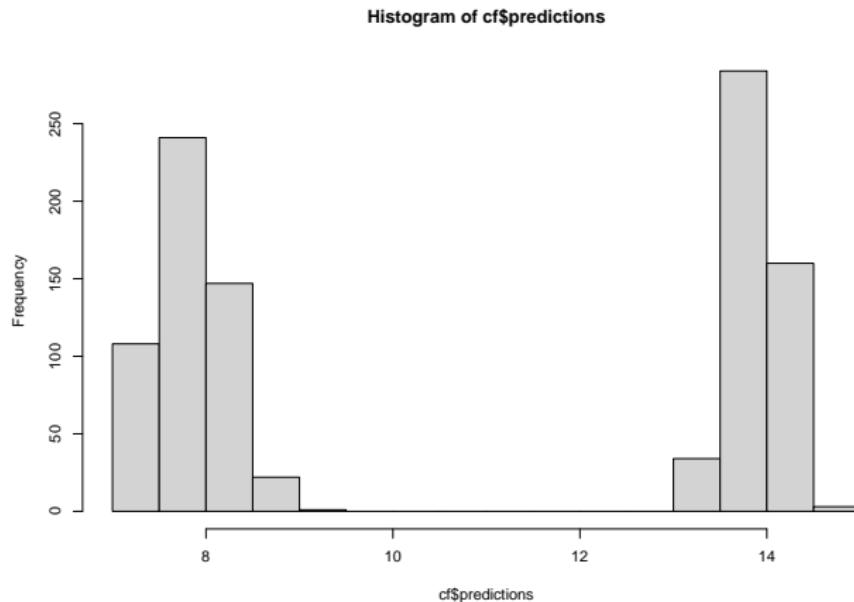
$$\mathbb{E}[\tau_i | X_i] = \mathbb{E}[Y_i(1) - Y_i(0) | X_i]$$

```
doordash$CATE = cf$predictions  
str(doordash)
```

```
## 'data.frame':    1000 obs. of  8 variables:  
##   $ ip        : num  10001 10002 10003 10004 10005 ...  
##   $ age       : num  59 63 59 39 27 22 45 21 43 40 ...  
##   $ mthsActive: num  19.26 11.51 4.26 9.53 19.54 ...  
##   $ treatment : num  0 0 1 1 0 1 1 1 1 0 ...  
##   $ mobile    : num  1 1 0 1 0 0 1 1 1 1 ...  
##   $ seconds   : num  116.3 146.6 159.8 120.9 40.3 ...  
##   $ purchase  : int  0 0 1 0 0 0 1 0 0 0 ...  
##   $ CATE      : num [1:1000, 1] 13.97 13.77 7.43 13.76 8
```

# Heterogeneity in predicted CATES

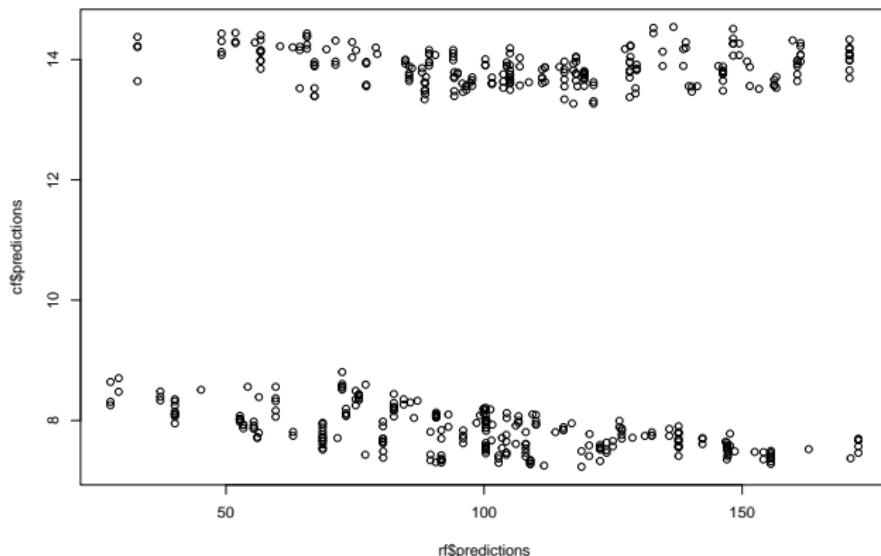
```
hist(cf$predictions)
```



## Causal forest versus random forest

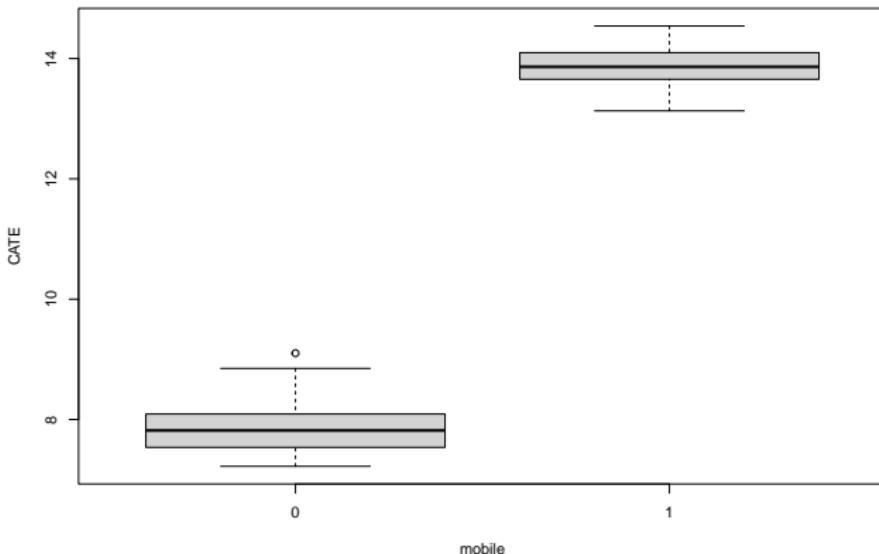
The random forest predicts *seconds* but the causal forest predicts *CATEs*

```
plot(rf$predictions, cf$predictions)
```



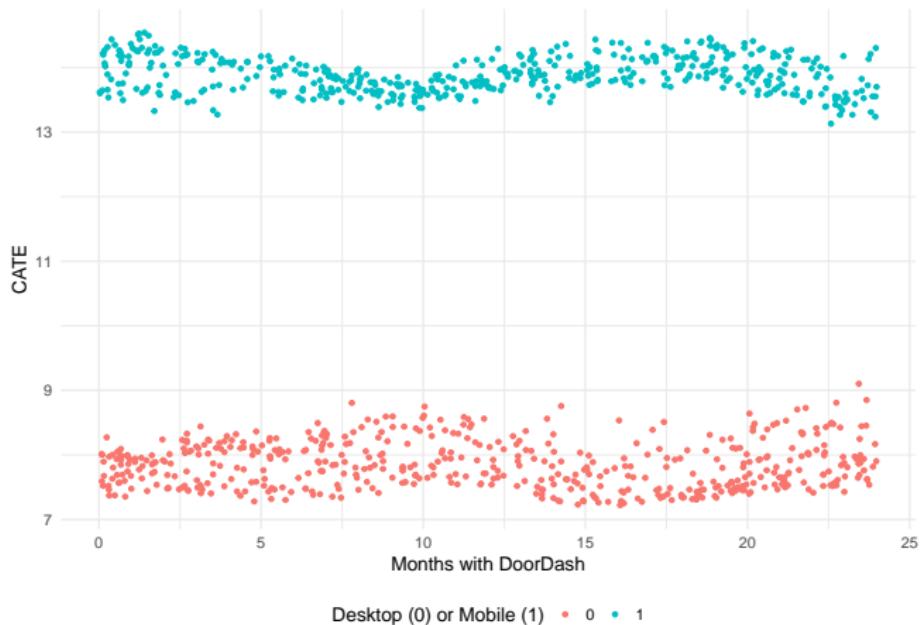
## Causal forest CATES by desktop/mobile

```
boxplot(CATE~mobile,doordash)
```



```
#or stripchart(CATE~mobile,doordash,vertical=TRUE)
```

## Causal forest CATEs versus age



## Optimal treatment given $X$ ?

```
library(policytree)
#opportunity cost of not treating = -tau
rewards = cbind(control=-get_scores(cf),
                 treatment=get_scores(cf))
tree = policy_tree(X, rewards, min.node.size = 1)
plot(tree, leaf.labels=c("Control", "Treatment"))
```

*#see R for this plot*

## Which predictors are the most important?

It can be hard to figure out which modifiers are “important”

One way to do that is to use `grf:best_linear_projection` which finds the linear model that fits best to the random forest.

```
best_linear_projection(cf, X)
```

```
##  
## Best linear projection of the conditional average treatment effect  
## Confidence intervals are cluster- and heteroskedasticity-robust  
##  
##           Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 8.529637   1.111182  7.6762 3.901e-14 ***  
## mobile      5.754905   0.584360  9.8482 < 2.2e-16 ***  
## age        -0.027341   0.024240 -1.1279   0.2596  
## mthsActive  0.033474   0.047492  0.7048   0.4811  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

## Causal forest average treatment effect

`grf::average_treatment_effect` combines the CATES to obtain an average treatment effect

It is similar to our difference-in-means estimate

```
average_treatment_effect(cf)
```

```
##   estimate    std.err  
## 10.576744  0.298686
```

Sometimes `std.err` of the causal forest ATE will be smaller than the difference-in-means ATE.

Accounting for variability due to differences in  $X_i$  can make the ATE estimate more precise.

## Power Calculations

# Hypothesis testing review

## Alpha $\alpha$

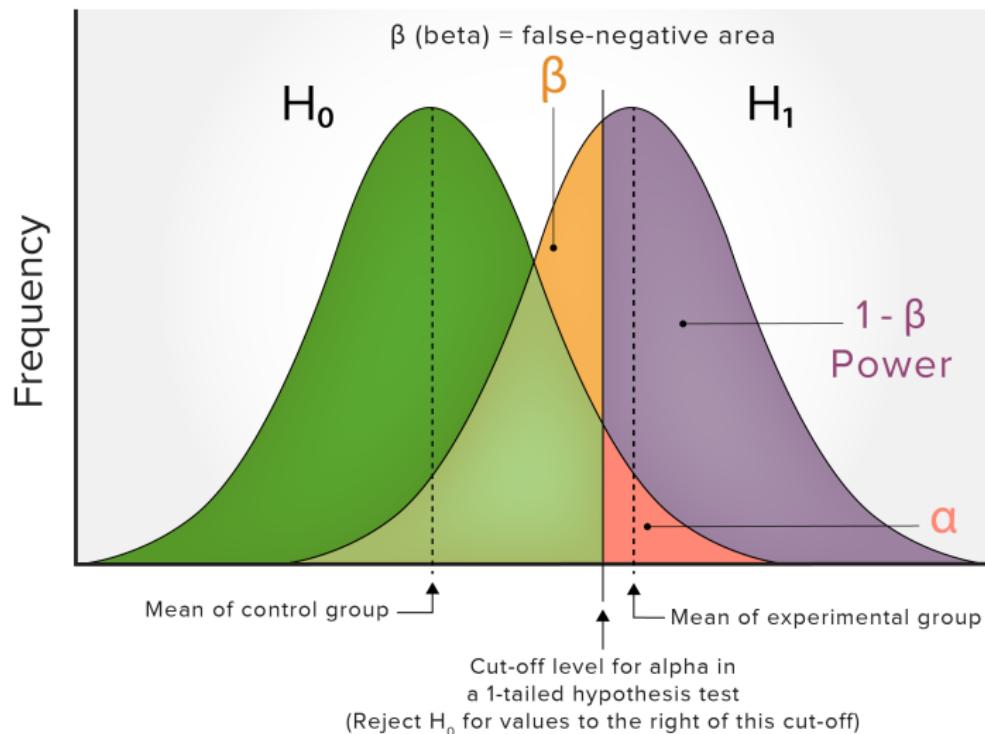
- ▶ Significance level (usually .05)
- ▶ Probability of making a Type I error
- ▶ A Type I error is a “false positive”
- ▶ False positive: an incorrect “yes”, or we think there *is* an effect when there *is not*
- ▶ Specifically: *rejecting* the null when the null is actually *true*

## Beta $\beta$

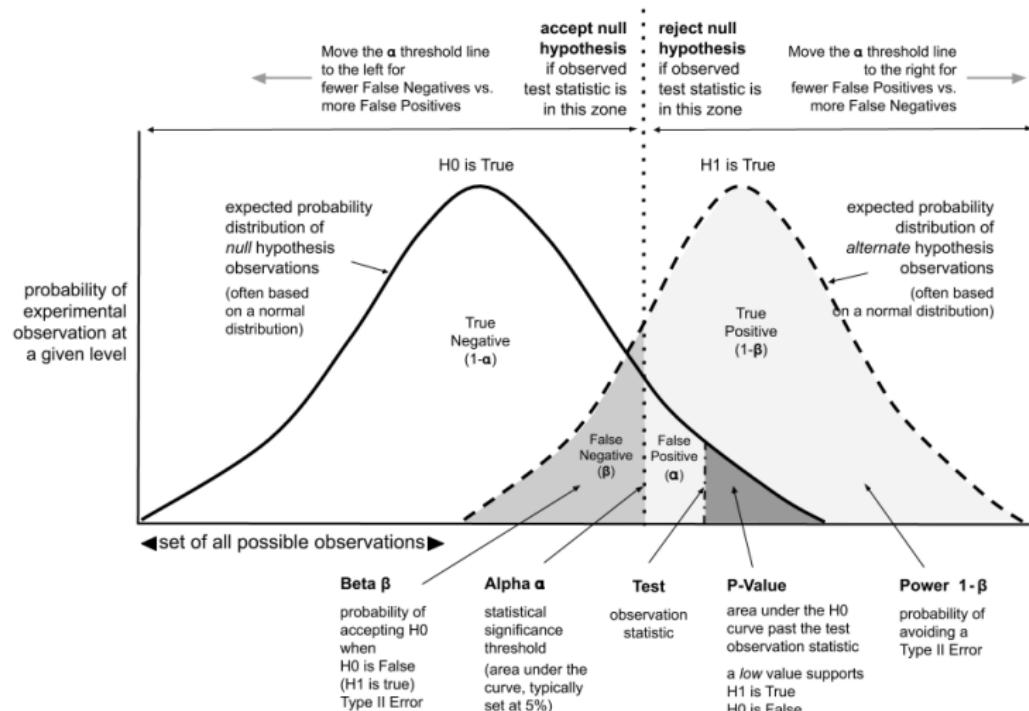
- ▶ Probability of a Type II error
- ▶ A Type II error is a “false negative”
- ▶ False negative: an incorrect “no”, or we think there *is no* effect when there *is*
- ▶ Specifically: *failing* to reject the null when the null is actually *false*

*Note:*  $\alpha \neq (1 - \beta)$ , but they are inversely related: increasing  $\alpha$  decreases  $\beta$

# One picture...



## Another option...



## The Importance of Power

**Power** is  $1 - \beta$ : the probability of finding a effect when a true effect exists

In other words, reject the null hypothesis when it is in fact false

Typically set to .8

How to increase power:

- ▶ Increase the sample size
- ▶ Reduce variance (noise in the data) by controlling for confounding variables and make units as similar as possible
- ▶ Increase alpha (say .05 to .10), but this increases the probability of a Type I error (false positive)

# Power Calculations in R

See library(pwr)

```
pwr.t.test(d = 0.2,           #difference in means
            sig.level = 0.05,      #Type I error (alpha)
            power = 0.8,          #Power (1 - pr(Type II))
            type = "two.sample",  #one/two/paired
            alternative = "two.sided")
```

```
##  
##      Two-sample t test power calculation  
##  
##              n = 393.4057  
##              d = 0.2  
##      sig.level = 0.05  
##      power = 0.8  
##      alternative = two.sided  
##  
## NOTE: n is number in *each* group
```

## Interpreting power calculations

The output shows the required sample size per group to achieve the desired power

A larger effect size requires a smaller sample size, and vice-versa

Next week we look at how to reduce variance to increase the power

## Smaller Effect → Larger N

```
pwr.t.test(d = 0.01,      #difference in means
            sig.level = 0.05,    #Type I error (alpha)
            power = 0.8,         #Power (1 - pr(Type II))
            type = "two.sample", #one/two/paired
            alternative = "two.sided")
```

```
##  
##      Two-sample t test power calculation  
##  
##              n = 156978.2  
##              d = 0.01  
##      sig.level = 0.05  
##      power = 0.8  
##      alternative = two.sided  
##  
## NOTE: n is number in *each* group
```

## Case Study: Amazon Website Optimization

# Goal

Increase conversion rate

- ▶ Percentage of visitors who make a purchase

**Hypothesis:** a redesigned product page with larger images and customer testimonials will lead to higher conversions

## Experimental Design

RCT: randomly assign website visitors to either the existing product page (control) or the redesigned page (treatment)

Metric: conversion rate (% of customers who end up purchasing)

Test duration: 2 weeks

## Data

```
load('data/amazon.rdata')
head(amazon)
```

	conversion	treatment	mobile	prime	monthlySpend
## 1	0	0	0	0	4.41
## 2	0	0	1	1	37.68
## 3	0	0	0	0	3.75
## 4	0	1	0	1	65.35
## 5	0	0	1	1	61.42
## 6	0	0	0	0	5.83

## ATE for proportions

```
ate_p = amazon %>%
  group_by(treatment) %>%
  summarise(conversion_bar = mean(conversion))

ate_p
```

```
## # A tibble: 2 x 2
##   treatment conversion_bar
##       <dbl>          <dbl>
## 1 0           0.0984
## 2 1           0.126
```

So the ATE is about 2.8%.

## Another way...

```
with(amazon,
      mean(conversion[treatment == 1]) -
      mean(conversion[treatment == 0]))  
  
## [1] 0.0277879
```

Is it significant?

## Confidence Intervals for Proportions

```
n = nrow(amazon)
binom.confint(sum(subset(amazon,treatment == 1)$conversion)
             n = sum(amazon$treatment),method='prop.test')

##           method     x     n      mean      lower      upper
## 1 prop.test 314 2489 0.1261551 0.113493 0.1399871

binom.confint(sum(subset(amazon,treatment == 0)$conversion)
             n = sum(amazon$treatment == 0),method='prop.t')

##           method     x     n      mean      lower      upper
## 1 prop.test 247 2511 0.09836718 0.08713672 0.1108456
```

## Use a prop-test

Testing the difference in proportions (i.e., purchase rate) between the treated and control groups

```
propData = amazon %>% group_by(treatment) %>%
  summarise(conversions = sum(conversion), n = n())
prop.test(propData$conversions, propData$n)
```

```
##  
## 2-sample test for equality of proportions with continu  
##  
## data: propData$conversions out of propData$n  
## X-squared = 9.4126, df = 1, p-value = 0.002155  
## alternative hypothesis: two.sided  
## 95 percent confidence interval:  
## -0.045675782 -0.009900014  
## sample estimates:  
##      prop 1      prop 2  
## 0.09836718 0.12615508
```

## Controlling for confounders

```
cate_p = glm(conversion ~ treatment*prime +
  monthlySpend + mobile,amazon,family='binomial')
#coef(cate_p)
tidy(cate_p) %>% select(1,2,5) %>%
  mutate(across(-1, ~round(., 3)))
```

```
## # A tibble: 6 x 3
##   term           estimate p.value
##   <chr>         <dbl>    <dbl>
## 1 (Intercept) -2.55      0
## 2 treatment     0.715     0
## 3 prime        0.075     0.579
## 4 monthlySpend 0.007     0
## 5 mobile        0.07      0.443
## 6 treatment:prime -1.02     0
```

## Confidence intervals

But these are on log-odds effects...

```
confint(cate_p)
```

```
## Waiting for profiling to be done...
```

```
##                                     2.5 %      97.5 %
## (Intercept)      -2.776116719 -2.335524123
## treatment        0.476479550  0.956953153
## prime            -0.190293374  0.340833209
## monthlySpend    0.005177537  0.008603113
## mobile           -0.108817882  0.249322626
## treatment:prime -1.394691840 -0.654120440
```

## Predict probabilities

Remember, treatment effect depends on  $X$  since this is non-linear

Setup various conditions and compare predicted conversion probabilities

```
cate_p_df = expand.grid(  
    treatment = 0:1,  
    mobile = 0:1,  
    monthlySpend = round(mean(amazon$monthlySpend), 2),  
    prime = 0:1)  
lo_hat = predict(cate_p, newdata = cate_p_df,  
                se.fit = TRUE)  
lower_lo = lo_hat$fit - 1.96 * lo_hat$se.fit  
upper_lo = lo_hat$fit + 1.96 * lo_hat$se.fit
```

## Compare results

Looks like Prime users have a null/negative treatment effect

```
cbind(cate_p_df,lb = round(plogis(lower_lo),3),  
      ub = round(plogis(upper_lo),3))
```

##	treatment	mobile	monthlySpend	prime	lb	ub
## 1	0	0	32.62	0	0.073	0.107
## 2	1	0	32.62	0	0.143	0.192
## 3	0	1	32.62	0	0.078	0.114
## 4	1	1	32.62	0	0.153	0.202
## 5	0	0	32.62	1	0.079	0.115
## 6	1	0	32.62	1	0.058	0.089
## 7	0	1	32.62	1	0.084	0.121
## 8	1	1	32.62	1	0.062	0.094

## \*Bootstrap SEs on the probabilities

Another option is to bootstrap the predicted probabilities

Probably more useful in more complex models, but might be useful to see here

Steps:

- 1) Sample with replacement from the data
- 2) Estimate the model and save predicted probabilities
- 3) Repeat many times

## \*In R

Using a simplified model, but the idea generalizes

```
nBootstraps = 100
n = nrow(amazon)
bootstrapOut = data.frame(control = rep(NA,nBootstraps),
                           treatment = rep(NA,nBootstraps))

for(b in 1:nBootstraps){
  df_b = amazon[sample(1:n,n,replace=TRUE),]
  temp = glm(conversion ~ treatment,family='binomial',data=df_b)
  bootstrapOut[b,] = predict(temp,data.frame(treatment = 0))
}
```

\*Now take quantiles

```
#head(bootstrapOut)
apply(bootstrapOut, 2, quantile)

##           control treatment
## 0%    0.08320189 0.1092808
## 25%   0.09384953 0.1233001
## 50%   0.09828241 0.1272361
## 75%   0.10194856 0.1309408
## 100%  0.11353712 0.1426269
```

\*Or get quantiles this way...

```
bootstrapOut %>%
  reframe(across(c(control,treatment),
                 ~quantile(.x,c(.025,.975))))
```

```
##           control treatment
## 1  0.08767708  0.1112726
## 2  0.10940749  0.1394809
```

## \*Another option for bootstrapping...

Slight modification

```
singleBootstrap = function() {  
  df_b = amazon[sample(1:n,n,replace=TRUE),]  
  temp = glm(conversion ~ treatment,family='binomial',data=df_b)  
  predict(temp,data.frame(treatment = 0:1),type='response')  
}  
  
#control / treated  
bootOut = t(replicate(10,singleBootstrap()))  
apply(bootOut,2,quantile)  
  
## 1 2  
## 0% 0.08660194 0.1133936  
## 25% 0.08874525 0.1238639  
## 50% 0.09371650 0.1272913  
## 75% 0.09646334 0.1307292  
## 100% 0.10760741 0.1384430
```

## Causal forest

```
W = amazon$treatment
Y = amazon$conversion
X = amazon %>% select(-c(treatment,conversion))
cf_amzn = causal_forest(X, Y, W.hat = 0.5, seed=1)
amazon$CATE = cf_amzn$predictions
```

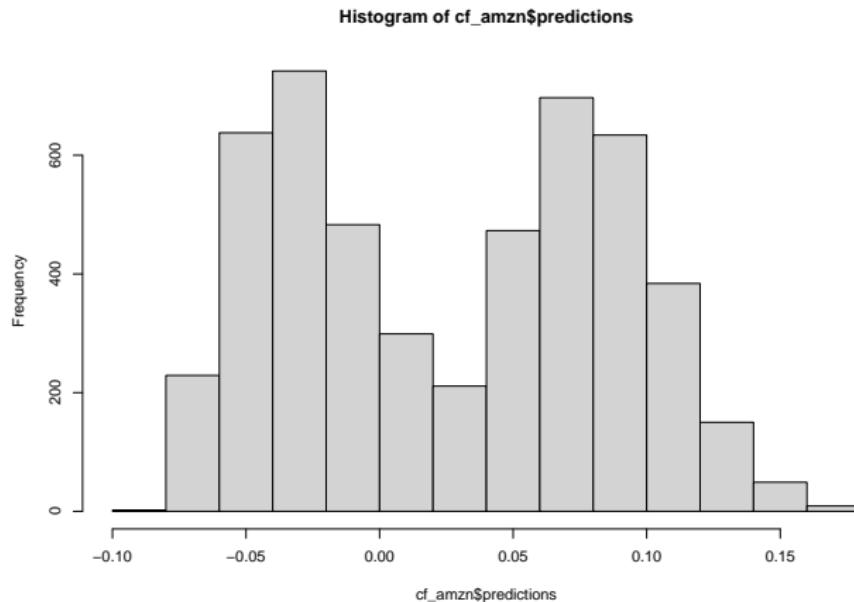
## Causal forest ATE

```
average_treatment_effect(cf_amzn)
```

```
##      estimate      std.err
## 0.027855967 0.008837823
```

# Heterogeneity in predicted CATES

```
hist(cf_amzn$predictions)
```



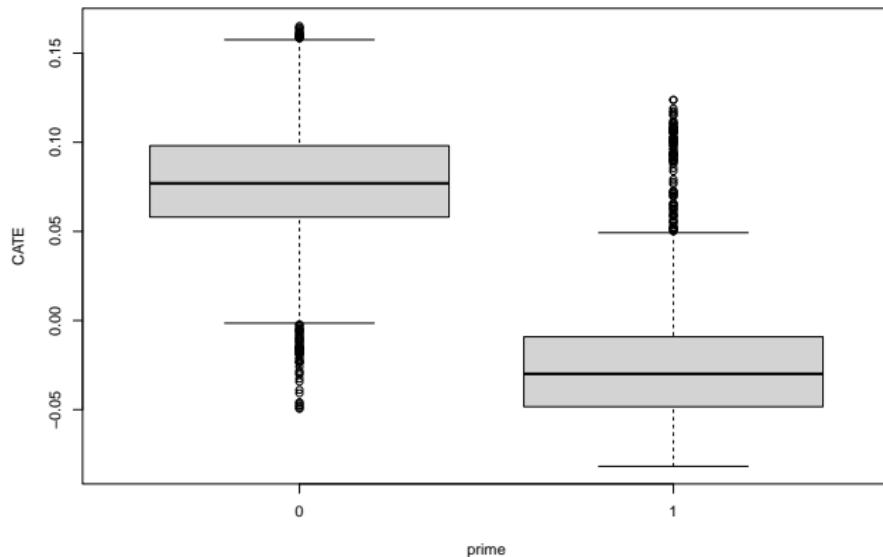
## Which predictors are the most important?

```
best_linear_projection(cf_amzn, X)
```

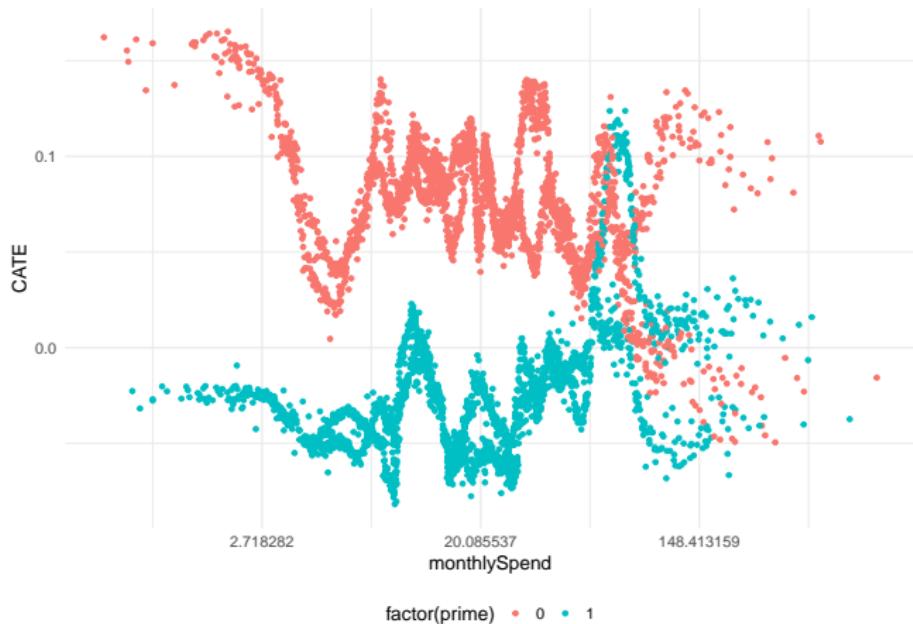
```
##  
## Best linear projection of the conditional average treatment effect  
## Confidence intervals are cluster- and heteroskedasticity-robust  
##  
##           Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 7.8249e-02 1.7333e-02 4.5143 6.497e-06 ***  
## mobile      7.1840e-03 1.7590e-02 0.4084  0.6830  
## prime       -1.0507e-01 1.7611e-02 -5.9663 2.594e-09 ***  
## monthlySpend -6.9163e-05 2.9865e-04 -0.2316  0.8169  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

## Causal forest CATES versus mobile users

```
boxplot(CATE~prime,amazon)
```



## Causal forest CATEs versus Prime



## Which treatment given $X$ ?

```
rewards = cbind(control=-get_scores(cf_amzn),  
                 treatment=get_scores(cf_amzn))  
tree = policy_tree(X, rewards, min.node.size = 1)  
plot(tree, leaf.labels=c("Control", "Treatment"))
```

## Interpretation

- ▶ On average, the redesigned page increased conversions by about 2.8%
- ▶ But, the treatment effect was significantly lower for Prime users (why? Maybe they were so used to the old layout)
- ▶ We also see heterogeneity in monthly spend levels: the Prime differential is greater for those with lower monthly spend

## What to do?

What are the limitations of releasing this?

Is the negative Prime effect a concern?

Is the difference worth it?

- ▶ Think about proportion of Prime users to everyone else, if it is small then it doesn't matter if there is a negative effect in this group

Any other ideas?

## Bonus Tips

## On Writing I of II

- ▶ **For your analysis to have impact, it needs to be understood**
- ▶ Write like you speak: plainly with small words, simple sentences, and short paragraphs
- ▶ Read your writing out loud. Good writing sounds natural.
- ▶ Be specific: “increases substantially” → “increases by 5%”
- ▶ Focus on actionable insights

## Writing Tips II

Use active voice rather than passive voice:

{subject} + {verb} + {object}

### **Active**

- ▶ Our campaign generated 15% more leads
- ▶ We collected survey responses from 5,000 subjects
- ▶ We recommend implementing an A/B test on the home page

### **Passive**

- ▶ 15% more leads were generated from our campaign
- ▶ Survey responses were collected from 5,000 subjects
- ▶ It is recommended that A/B testing is implemented on the home page

## Next Week

- ▶ Variance reduction with CUPED
- ▶ PSM
- ▶ AIPW
- ▶ DoubleML